

California Housing Price Prediction

A comprehensive regression analysis project built using Python and **Streamlit** to predict median housing prices in California. This application allows users to explore data, train models with flexible configurations, evaluate model performance, and make predictions — all through an interactive web interface.

Table of Contents

1. [Project Objectives](#)
2. [Dataset Description](#)
3. [Data Preprocessing & Feature Engineering](#)
4. [Application Architecture](#)
5. [Model Training & Evaluation](#)
6. [Prediction Module](#)
7. [Results Summary](#)
8. [Challenges & Solutions](#)
9. [Future Enhancements](#)
10. [How to Run the App](#)
11. [Dependencies](#)

Project Objectives

The main goals of this project are:

- Predict the **Median House Value** in California using demographic and geographic features.
- Implement **end-to-end regression modeling**, including EDA, feature engineering, model building, and evaluation.
- Provide a **user-friendly interface** for training, evaluating, and deploying models interactively using Streamlit.

Dataset Description

The dataset is sourced from `sklearn.datasets.fetch_california_housing` and includes:

- **Instances:** 20,640

- **Features** (original):
 - MedInc: Median income in block
 - HouseAge: Median house age
 - AveRooms: Average number of rooms
 - AveBedrms: Average number of bedrooms
 - Population: Block population
 - AveOccup: Average occupants per household
 - Latitude, Longitude: Geographical coordinates
- **Target:**
 - MedHouseVal: Median house value in \$100,000s

Data Preprocessing & Feature Engineering

Cleaning

- Checked for and confirmed the absence of **missing values**.
- Applied **outlier capping** (winsorization) to reduce the effect of extreme values using the 1st and 99th percentiles.

Feature Engineering

- Added multiple domain-driven features:
 - $\text{RoomsPerPopulation} = \text{AveRooms} * \text{AveOccup}$
 - $\text{BedroomsRatio} = \text{AveBedrms} / \text{AveRooms}$
 - $\text{IncomePerRoom} = \text{MedInc} / \text{AveRooms}$
 - $\text{LatxLong} = \text{Latitude} * \text{Longitude}$ (interaction term to reflect spatial effects)

Normalization

- Standardized features using StandardScaler.

Application Architecture

Built using **Streamlit**, the app is divided into four main pages:

1. Data Exploration

- Descriptive statistics

- Missing values check
- Correlation matrix (heatmap)
- Distribution plots (histograms, boxplots)
- Geospatial plot (colored scatter plot with population-based size)

2. Model Training & Evaluation

- **Model options:**
 - Linear Regression
 - Ridge, Lasso, ElasticNet
 - Polynomial Regression (degree=2)
- **Feature Selection:** via `SelectKBest(f_regression)`
- **Dimensionality Reduction:** Optional PCA
- **Cross-Validation:** Customizable folds
- **Metrics Calculated:**
 - MSE, RMSE, MAE, R^2 (for both train/test)
- **Model Analysis:**
 - Coefficient ranking
 - Feature importance visualization
 - Residual plot & Q-Q plot
- **Model Export:** Trained models are saved using `joblib`.

3. Price Prediction

- Dynamic form to input feature values
- Reconstructs engineered features automatically
- Shows predicted price in both \$100,000s and \$ units
- (If supported by model): Visualizes **feature contributions**

4. Project Report

- Presents project overview, findings, and technical notes

Model Training & Evaluation

Modeling Pipeline:

1. **Preprocessing Steps:**
 - a. Scaling

- b. Optional PCA
 - c. Optional K-Best feature selection
- 2. **Regression Model** (selectable)
- 3. **Optional PolynomialFeatures transformation**

Regularization Support:

- Ridge, Lasso, and ElasticNet allow tuning of α via Streamlit sliders.

Cross-Validation:

- Implemented using `cross_val_score()` from Scikit-learn.
- Configurable number of folds (3 to 10).

Evaluation Metrics:

Metric	Description
MSE	Mean Squared Error
RMSE	Root Mean Squared Error
MAE	Mean Absolute Error
R^2	Coefficient of Determination

Prediction Module

- Loads the trained model from disk.
- Takes input features from a web form.
- Automatically computes engineered features.
- Outputs:
 - Predicted median house value in both 100k and dollar units.
 - If model supports `.coef_`, shows:
 - Contribution of each feature
 - Bar plot of top contributors

Results Summary

- **Best Performing Model:** Ridge Regression
- **Test R^2 Score:** ~0.62
- **Top Predictive Features:**
 - MedInc (median income)

- Latitude, Longitude
- AveRooms
- RoomsPerPopulation
- **Model saved** as `housing_model.pkl`

Challenges & Solutions

Challenge	Solution
Multicollinearity (AveRooms, AveBedrms)	Used Lasso/Ridge and SelectKBest
Outliers skewing the distribution	Winsorization (1st, 99th percentile)
Non-linearity in features	Polynomial regression + PCA
Model interpretability in high-dim. space	Restricted polynomial degree to 2 and visualized top features

Future Enhancements

- Integrate **XGBoost**, **Random Forest**, or **Gradient Boosting**
- Perform **hyperparameter tuning** using GridSearchCV/Optuna
- Add **real-time map visualizations** with Folium or Plotly
- Enable **model comparison dashboards**
- Convert to **full-stack app** using Flask + Streamlit front-end

Dependencies

```

numpy
pandas
matplotlib
seaborn
scikit-learn

statsmodels
joblib
streamlit
pillow

```

